

# Computer System Operation

# System Startup

- On power up
  - everything in system is in random, unpredictable state
  - special hardware circuit raises RESET pin of CPU
    - sets the program counter to 0xffffffff
    - this address is mapped to ROM (Read-Only Memory)
- BIOS (Basic Input/Output Stream)
  - set of programs stored in ROM
  - some OS's use only these programs
    - MS DOS
  - many modern systems use these programs to load other system programs
    - Windows, Unix, Linux

# BIOS

- General operations performed by BIOS
  - 1) find and test hardware devices
    - POST (Power-On Self-Test)
  - 2) initialize hardware devices
    - creates a table of installed devices
  - 3) find *boot sector*
    - may be on floppy, hard drive, or CD-ROM
  - 4) load boot sector into memory location 0x00007c00
  - 5) sets the program counter to 0x00007c00
    - starts executing code at that address

# Boot Loader

- Small program stored in boot sector
- Loaded by BIOS at location 0x00007c0
- Configure a basic file system to allow system to read from disk
- Loads kernel into memory
- Also loads another program that will begin kernel initialization

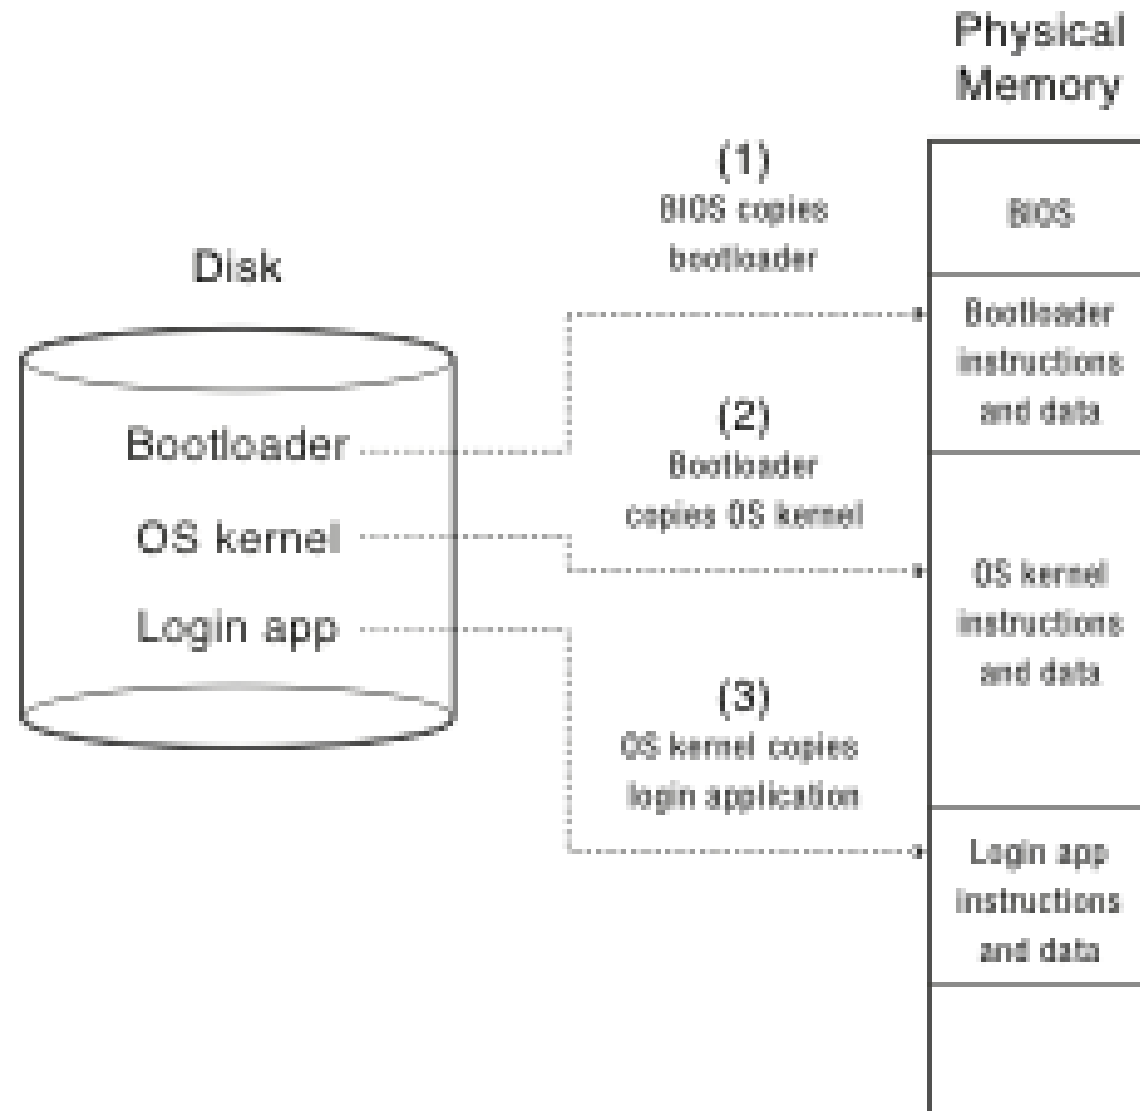
# Initial Kernel Program

- Determines amount of RAM in system
  - uses a BIOS function to do this
- Configures hardware devices
  - video card, mouse, disks, etc.
  - BIOS may have done this but usually redo it
    - portability
- Switches the CPU from *real* to *protected* mode
  - real mode: fixed segment sizes, 1 MB memory addressing, and no segment protection
  - protected mode: variable segment sizes, 4 GB memory addressing, and provides segment protection
- Initializes paging (virtual memory)

# Final Kernel Initialization

- Sets up page tables and segment descriptor tables
  - these are used by virtual memory and segmentation hardware
- Sets up interrupt vector and enables interrupts
- Initializes all other kernel data structures ( Linked lists, Binary search trees, Bitmaps etc.)
- Creates initial process and starts it running
  - *init* in Linux
  - *smss* (Session Manager SubSystem) in NT

# Booting



# System Programs

- Application programs included with the OS
- Highly trusted programs
- Perform useful work that most users need
  - listing and deleting files, configuring system
  - ls, rm, Windows Explorer and Control Panel
  - may include compilers and text editors
- Not part of the OS
  - run in user space
- Very useful

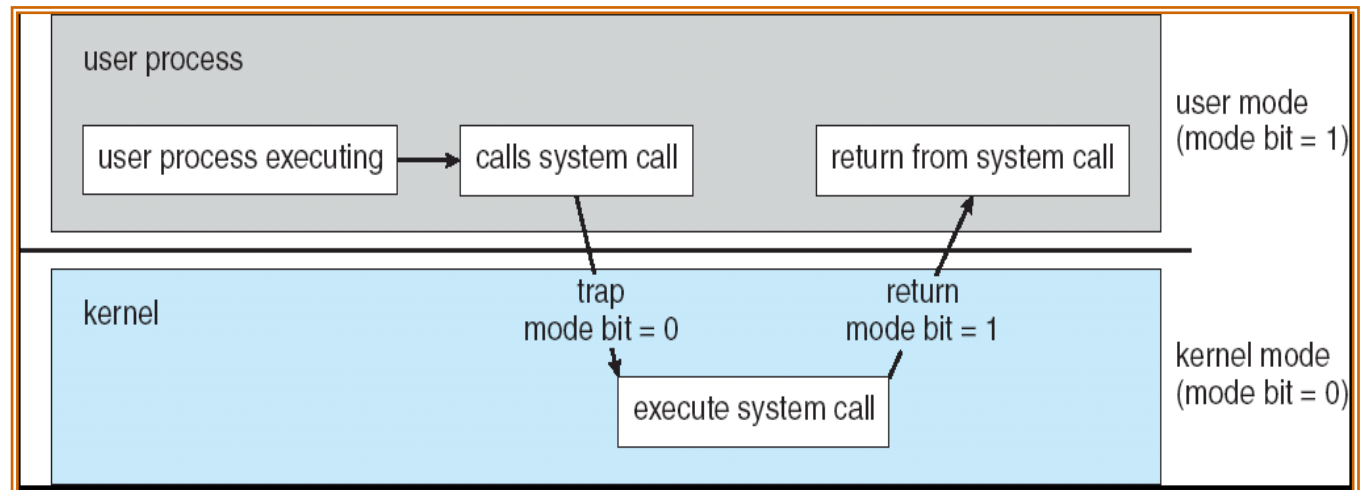


# Operating-System Operations

- Interrupt driven by hardware
- Software error or request creates **exception** or **trap**
  - Division by zero, request for operating system service
- Other process problems include infinite loop, processes modifying each other or the operating system
- **Dual-mode** operation allows OS to protect itself and other system components
- Hardware provides at least two modes:
  - “Kernel” mode (or “supervisor” or “protected”)
  - “User” mode: Normal programs executed
  - **Mode bit** provided by hardware
    - Provides ability to distinguish when system is running user code or kernel code
    - Some instructions designated as **privileged**, only executable in kernel mode
    - System call changes mode to kernel, return from call resets it to user

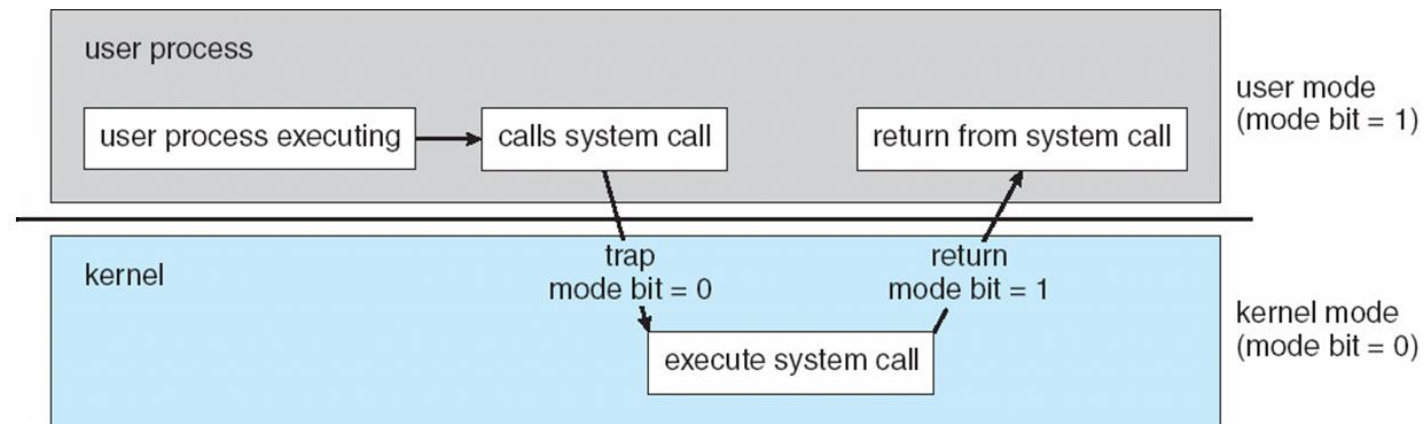
# Dual Mode Operation

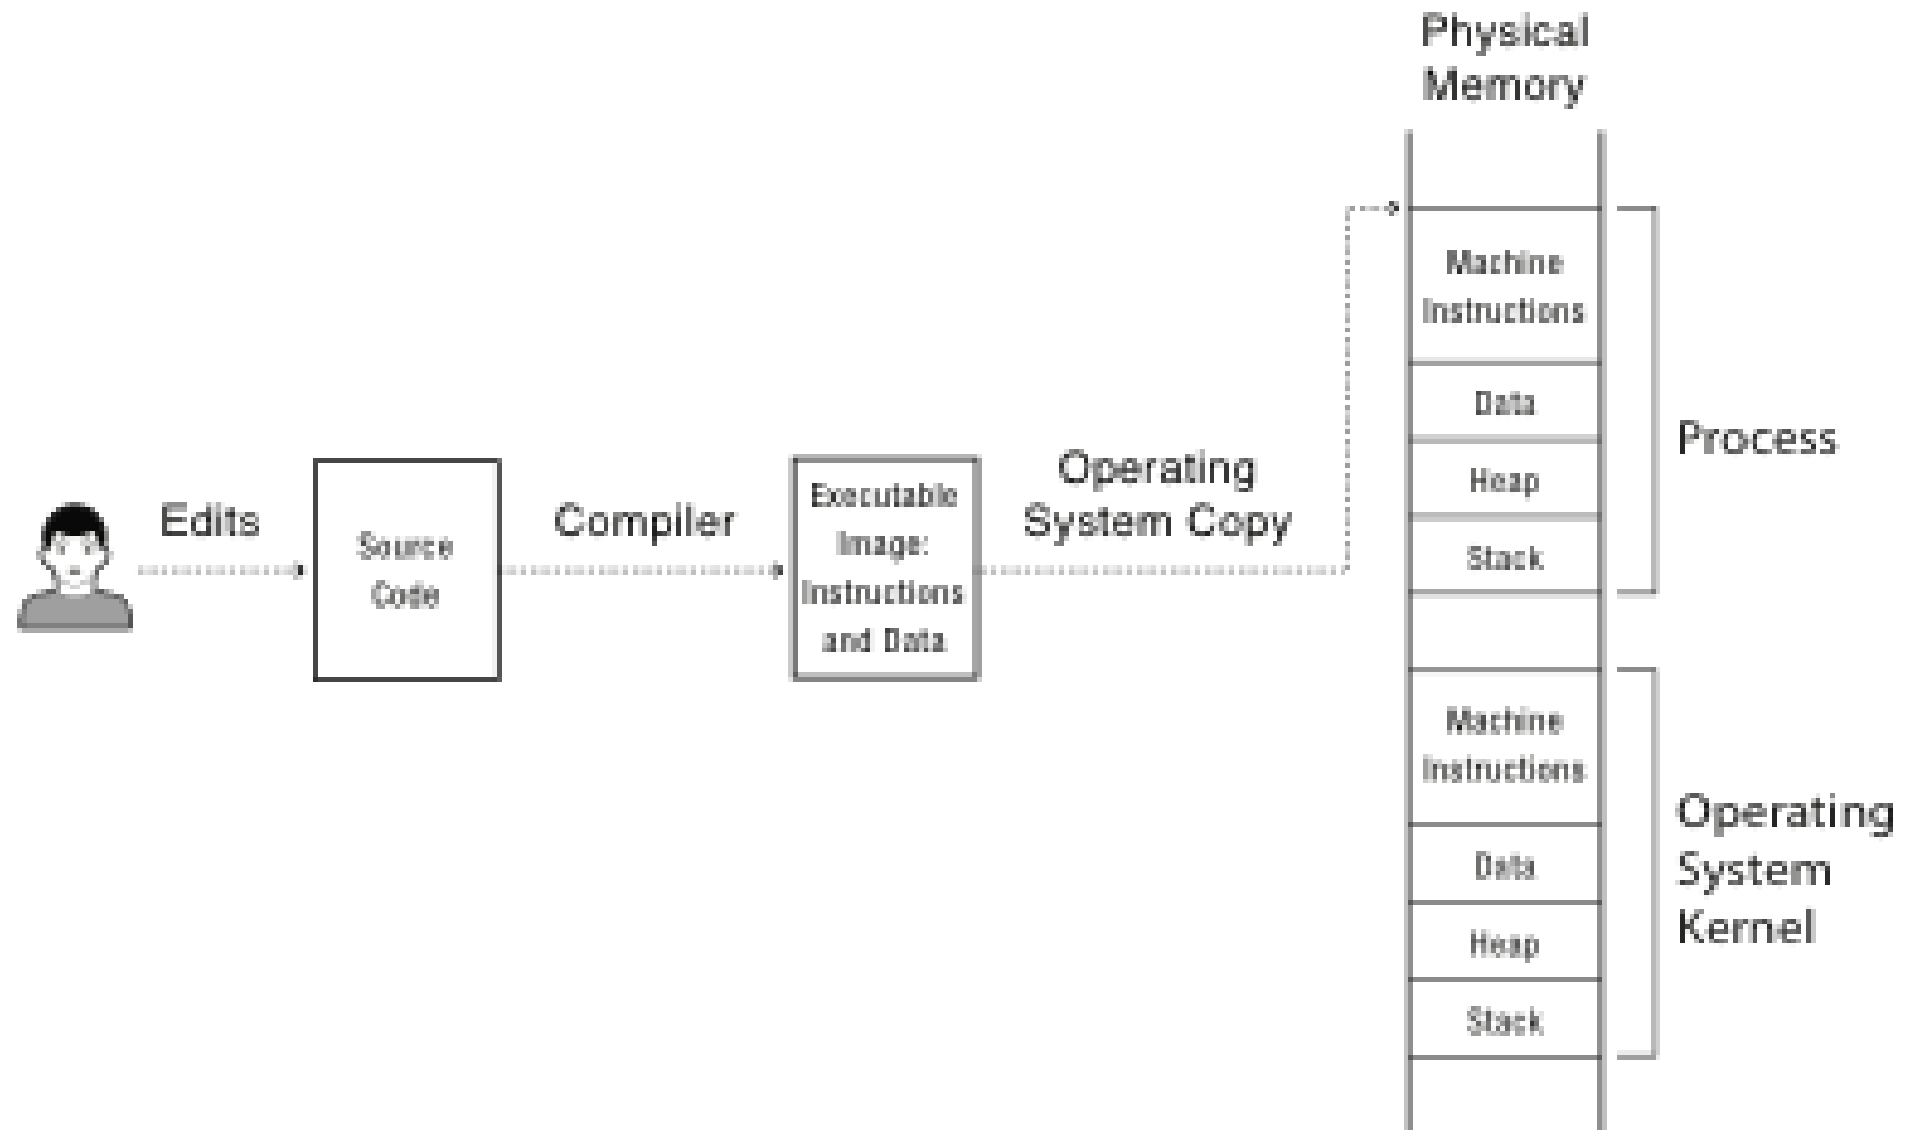
- Some instructions/ops prohibited in user mode:
  - Example: cannot modify page tables in user mode
    - Attempt to modify  $\Rightarrow$  Exception generated
- Transitions from user mode to kernel mode:
  - System Calls, Interrupts, Other exceptions



# Transition from User to Kernel Mode

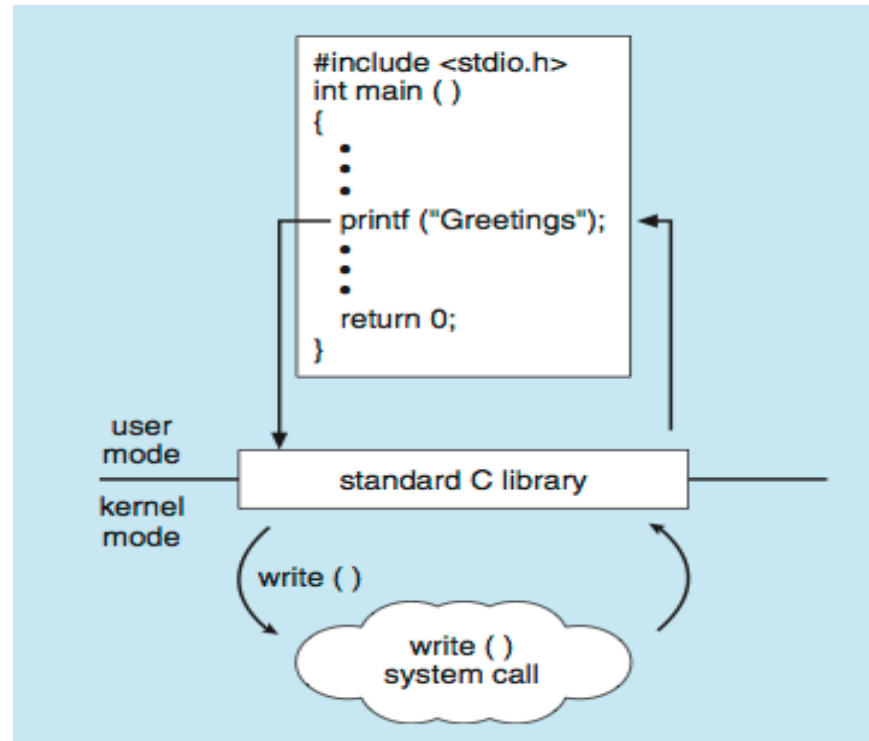
- Timer to prevent infinite loop / process hogging resources
  - Set interrupt after specific period
  - Operating system decrements counter
  - When counter zero generate an interrupt
  - Set up before scheduling process to regain control or terminate program that exceeds allotted time





# Standard C Library Example

- C program invoking printf() library call, which calls write() system call



# System Calls


- System calls provide the interface between a running program and the operating system.
  - Generally available as assembly-language instructions.
  - Languages have been defined to replace assembly language for systems programming; allow system calls to be made directly (e.g., C, C++)
- Three general methods are used to pass parameters between a running program and the operating system.
  - Pass parameters in *registers*.
  - Store the parameters in a table in memory, and the table address is passed as a parameter in a register.
  - *Push* (store) the parameters onto the *stack* by the program, and *pop* off the stack by operating system.

# System Calls

- System calls are routines run by the OS on behalf of the user
- Allow user to access I/O, create processes, get system information, etc.
- Programming interface to the services provided by the OS
- Typically written in a high-level language (C or C++)
- Mostly accessed by programs via a high-level **Application Programming Interface (API)** rather than direct system call use
- Three most common APIs are Win32 API for Windows, POSIX API for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X), and Java API for the Java virtual machine (JVM)

# System Calls

- A System Call is the main way a user program interacts with the Operating System.






# System Calls

## HOW A SYSTEM CALL WORKS

- Obtain access to system space
- Do parameter validation
- System resource collection ( locks on structures )
- Ask device/system for requested item
- Suspend waiting for device
- Interrupt makes this thread ready to run
- Wrap-up
- Return to user



There are 11 (or more) steps in making the system call  
**read (fd, buffer, nbytes)**

Linux API

# System Call Implementation

- Typically, a number associated with each system call
  - **System-call interface** maintains a table indexed according to these numbers
- The system call interface invokes the intended system call in OS kernel and returns status of the system call and any return values
- The caller need know nothing about how the system call is implemented
  - Just needs to obey API and understand what OS will do as a result call
  - Most details of OS interface hidden from programmer by API
    - Managed by run-time support library (set of functions built into libraries included with compiler)

# Example of Standard API

## EXAMPLE OF STANDARD API

As an example of a standard API, consider the `read()` function that is available in UNIX and Linux systems. The API for this function is obtained from the `man` page by invoking the command

```
man read
```

on the command line. A description of this API appears below:

<pre>#include &lt;unistd.h&gt;</pre>		
<pre>ssize_t</pre>	<pre>read(int fd, void *buf, size_t count)</pre>	
return	function	parameters
value	name	

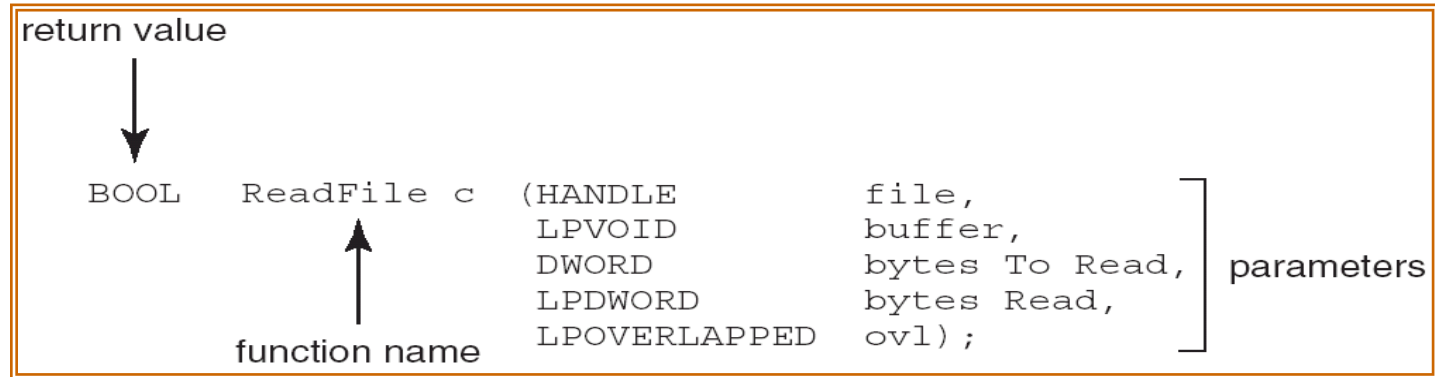
A program that uses the `read()` function must include the `unistd.h` header file, as this file defines the `ssize_t` and `size_t` data types (among other things). The parameters passed to `read()` are as follows:

- `int fd`—the file descriptor to be read
- `void *buf`—a buffer where the data will be read into
- `size_t count`—the maximum number of bytes to be read into the buffer

On a successful read, the number of bytes read is returned. A return value of 0 indicates end of file. If an error occurs, `read()` returns `-1`.

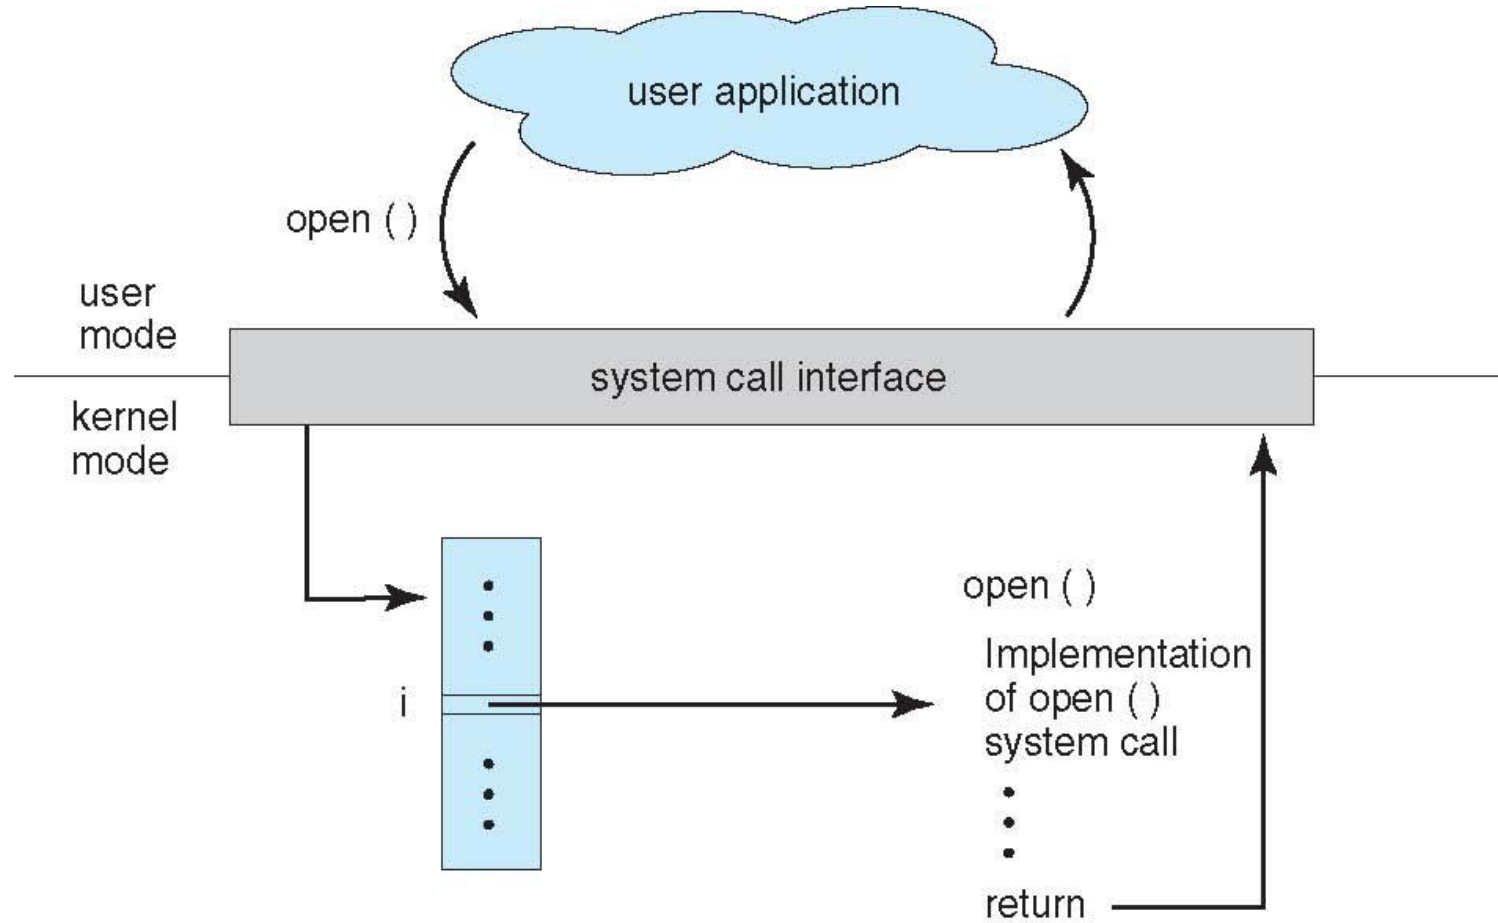
# Example of Windows API

- Consider the ReadFile() function in the program
- Win32 API—a function for reading from a file.



- A description of the parameters passed to `ReadFile()`
  - `HANDLE file`—the file to be read
  - `LPVOID buffer`—a buffer where the data will be read into and written from
  - `DWORD bytesToRead`—the number of bytes to be read into the buffer
  - `LPDWORD bytesRead`—the number of bytes read during the last read
  - `LPOVERLAPPED ovl`—indicates if overlapped I/O is being used


# API – System Call – OS Relationship



# System Call Parameter Passing

- Often, more information is required than simply identity of desired system call
  - Exact type and amount of information vary according to OS and call
- Three general methods used to pass parameters to the OS
  - Simplest: pass the parameters in registers
    - In some cases, may be more parameters than registers
  - Parameters stored in a block, or table, in memory, and address of block passed as a parameter in a register
    - This approach taken by Linux and Solaris
  - Parameters placed, or **pushed**, onto the **stack** by the program and **popped** off the stack by the operating system
  - Block and stack methods do not limit the number or length of parameters being passed

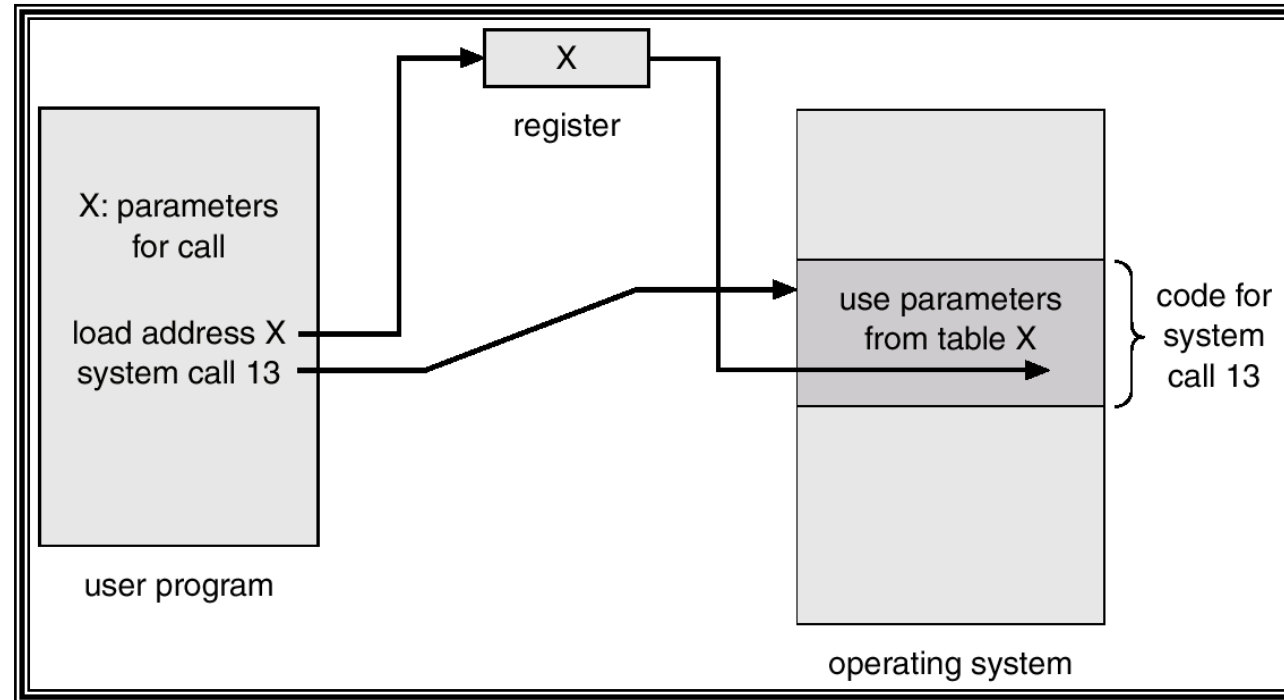
# passing data between programs



Msg Passing

Shared Memory

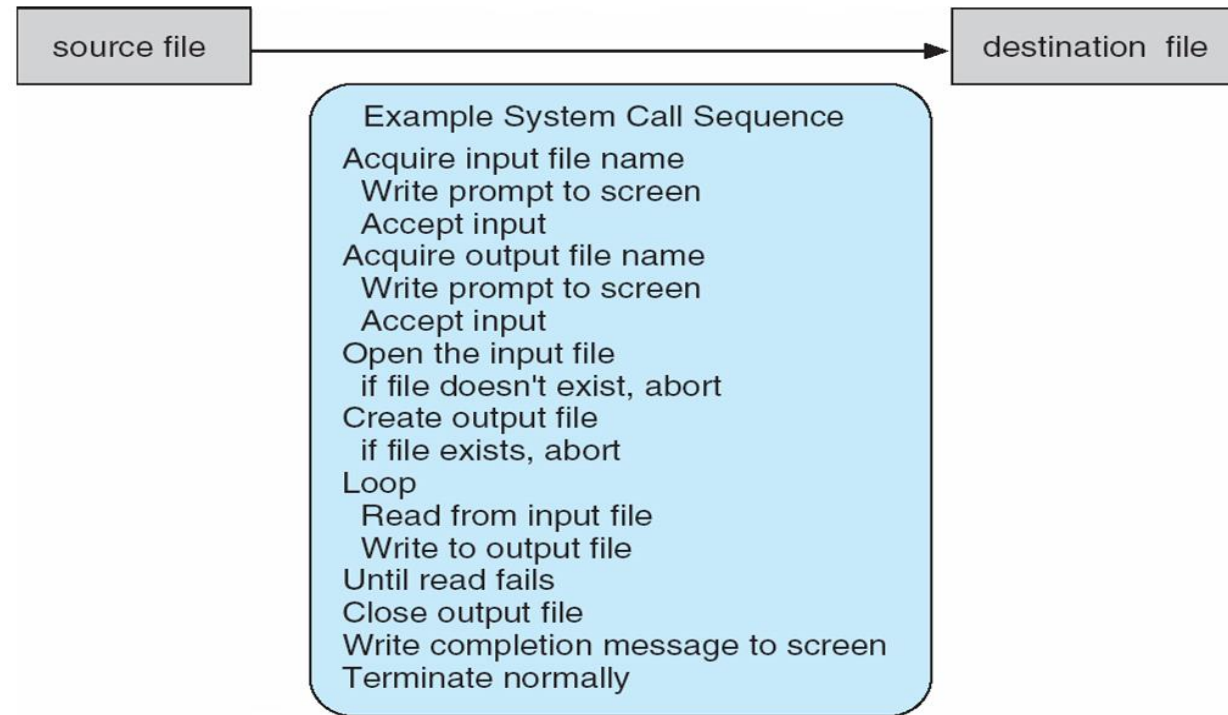
# Passing of Parameters As A Table





# Example of System Calls

- System call sequence to copy the contents of one file to another file



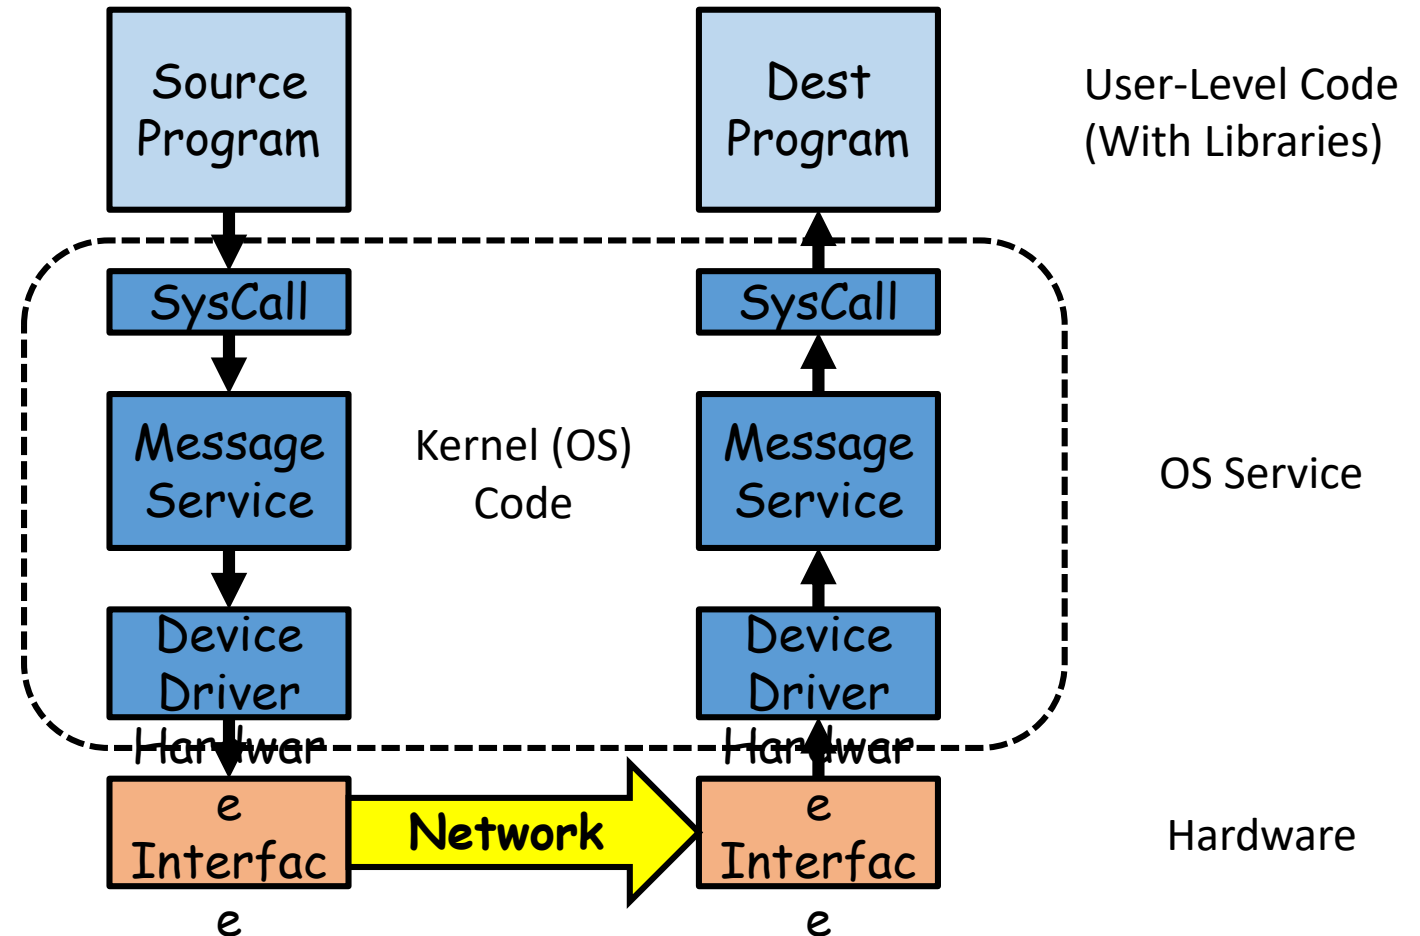
# Send message from one processor to another

Operations to be performed:

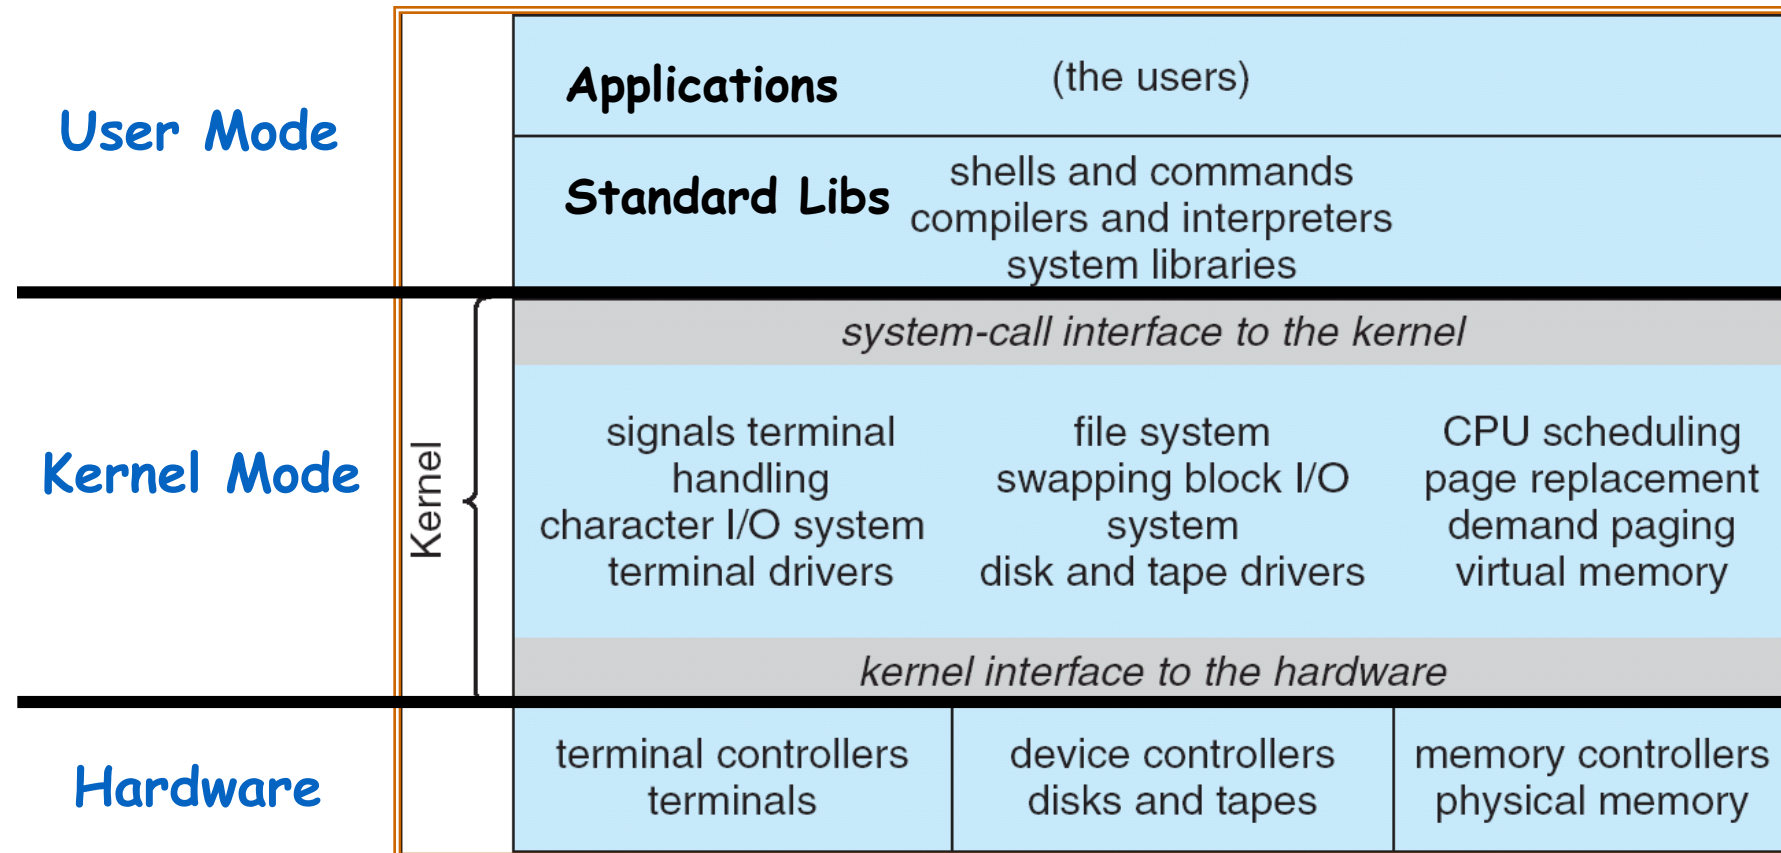
- Check Permissions, Format Message
- Enforce forward progress, Handle interrupts
- Prevent Denial Of Service (DOS) and/or Deadlock

Traditional Approach:

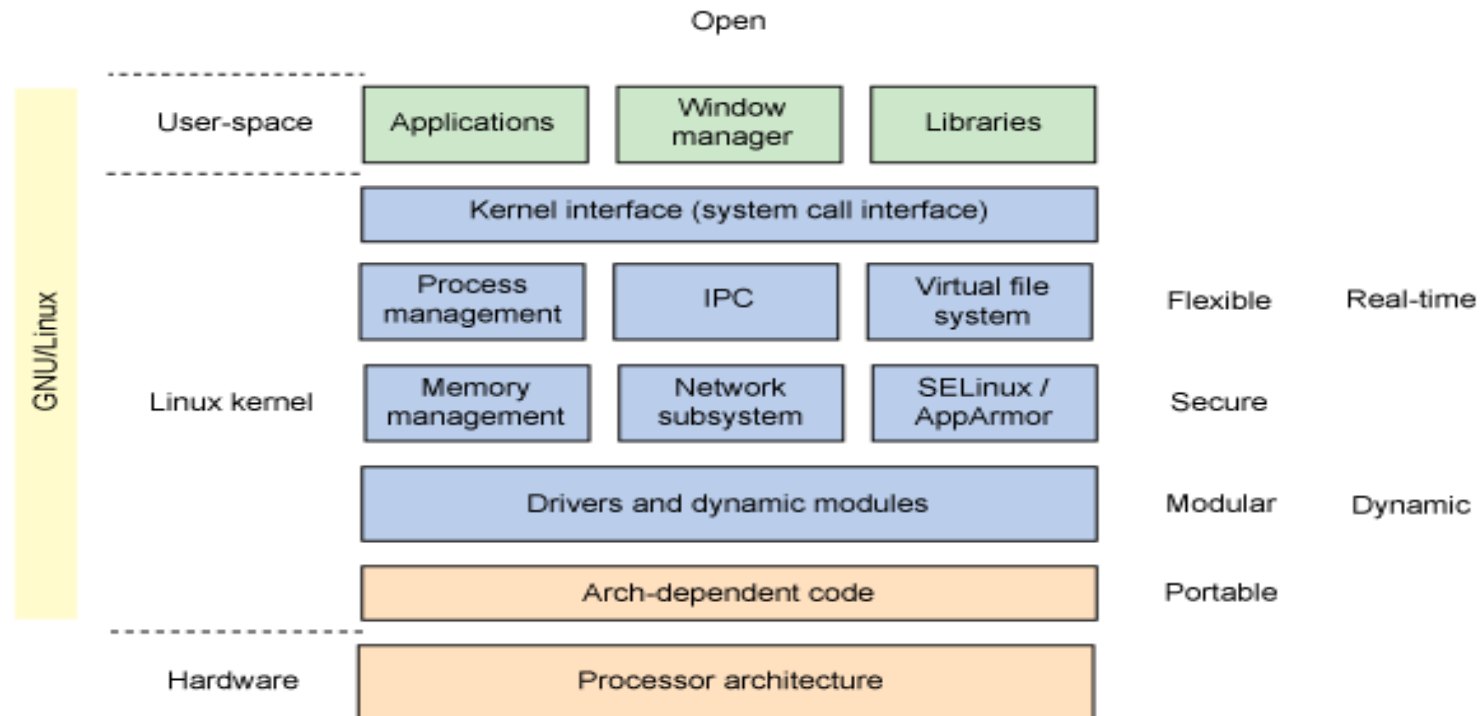
Use a system call + OS Service



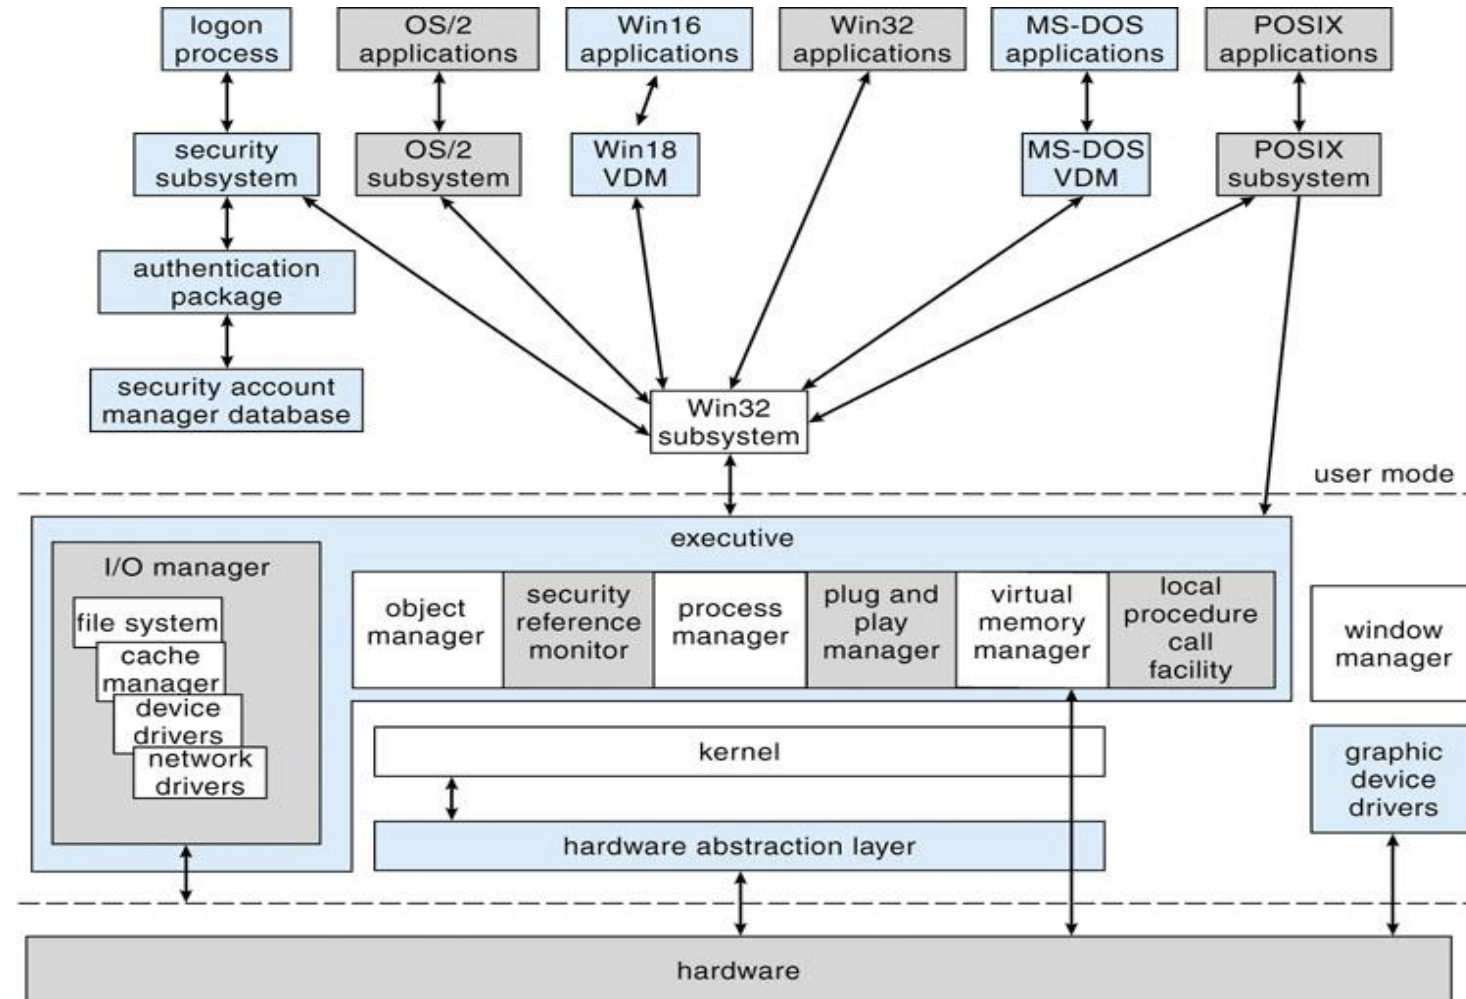
# UNIX System Structure



# Linux Structure



# Microsoft Windows Structure



# Major Windows Components

- Hardware Abstraction Layer
  - Hides hardware chipset differences from upper levels of OS
- Kernel Layer
  - Thread Scheduling
  - Low-Level Processors Synchronization
  - Interrupt/Exception Handling
  - Switching between User/Kernel Mode.
- Executive
  - Set of services that all environmental subsystems need
    - Object Manager
    - Virtual Memory Manager
    - Process Manager
    - Advanced Local Procedure Call Facility
    - I/O manager
    - Cache Manager
    - Security Reference Monitor
    - Plug-and-Plan and Power Managers
    - Registry
    - Booting
- Programmer Interface: Win32 API

# Types of System Calls

- Process control
- File manipulation
- Device manipulation
- Information maintenance
- Communications
- Protection

# Types of System Calls

- Process control
  - create process, terminate process
  - end, abort
  - load, execute
  - get process attributes, set process attributes
  - wait for time
  - wait event, signal event
  - allocate and free memory
  - Dump memory if error
  - **Debugger** for determining **bugs, single step** execution
  - **Locks** for managing access to shared data between processes



# Types of System Calls

- File management
  - create file, delete file
  - open, close file
  - read, write, reposition
  - get and set file attributes
- Device management
  - request device, release device
  - read, write, reposition
  - get device attributes, set device attributes
  - logically attach or detach devices

# Types of System Calls

- Information maintenance
  - get time or date, set time or date
  - get system data, set system data
  - get and set process, file, or device attributes
- Communications
  - create, delete communication connection
  - send, receive messages if **message passing model** to **host name** or **process name**
    - From **client** to **server**
  - **Shared-memory model** create and gain access to memory regions
  - transfer status information
  - attach and detach remote devices

# Types of System Calls

- Protection
  - Control access to resources
  - Get and set permissions
  - Allow and deny user access

# Examples of Windows and Unix System Calls

	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shmget() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()




<b>UNIX</b>	<b>Win32</b>	<b>Description</b>
fork	CreateProcess	Create a new process
waitpid	WaitForSingleObject	Can wait for a process to exit
execve	(none)	CreateProcess = fork + execve
exit	ExitProcess	Terminate execution
open	CreateFile	Create a file or open an existing file
close	CloseHandle	Close a file
read	ReadFile	Read data from a file
write	WriteFile	Write data to a file
lseek	SetFilePointer	Move the file pointer
stat	GetFileAttributesEx	Get various file attributes
mkdir	CreateDirectory	Create a new directory
rmdir	RemoveDirectory	Remove an empty directory
link	(none)	Win32 does not support links
unlink	DeleteFile	Destroy an existing file
mount	(none)	Win32 does not support mount
umount	(none)	Win32 does not support mount
chdir	SetCurrentDirectory	Change the current working directory
chmod	(none)	Win32 does not support security (although NT does)
kill	(none)	Win32 does not support signals
time	GetLocalTime	Get the current time

# Initial Hardware Model

The host machine is equipped with the physical hardware

e.g. a desktop with x-86 architecture running its installed Windows OS

All applications access hardware resources (i.e. memory, i/o) through system calls to operating system (privileged instructions)



- Single OS image per machine
- Software and hardware tightly coupled
- Running multiple applications on same machine
- Underutilized resources
- Inflexible and costly infrastructure


- Advantages
  - Design is decoupled (i.e. OS people can develop OS separate of Hardware people developing hardware)
  - Hardware and software can be upgraded without notifying the Application programs



- Disadvantages
  - Application compiled on one ISA will not run on another ISA..
    - Applications compiled for Mac use different operating system calls than application designed for windows.
  - ISA's must support old software
    - Can often be inhibiting in terms of performance
  - Since software is developed separately from hardware..
  - Software is not necessarily optimized for hardware.

# Architecture & Interfaces

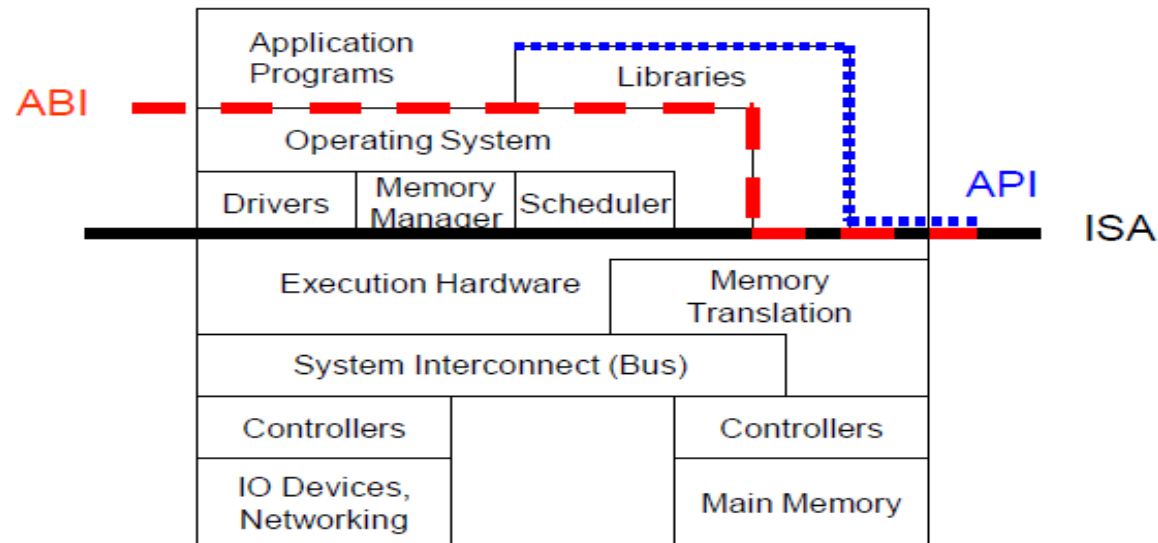
Architecture: formal specification of a system's interface and the logical behavior of its visible resources.



- **API** – Application Programming Interface
- **ABI** – Application Binary Interface
- **ISA** – Instruction Set Architecture

# Architecture, Implementation Layers

- Architecture
  - Functionality and Appearance of a computer system but not implementation details
  - Level of Abstraction = Implementation layer ISA, ABI, API



# Architecture, Implementation Layers

- Implementation Layer : ISA
  - Instruction Set Architecture
  - Divides hardware and software
  - Concept of ISA originates from IBM 360
    - Various prices, processing power, processing unit, devices
    - But guarantee a *software compatibility*
  - User ISA and System ISA

# Architecture, Implementation Layers

- Implementation Layer : ABI
  - Application Binary Interface
  - Provides a program with access to the hardware resource and services available in a system
  - Consists of User ISA and System Call Interfaces

# Architecture, Implementation Layers

- Implementation Layer : API
  - Application Programming Interface
  - Key element is Standard Library ( or Libraries )
  - Typically defined at the source code level of High Level Language
  - **clib** in Unix environment : supports the UNIX/C programming language

# Multimode System

- The concept of modes of operation in operating system can be extended beyond the dual mode. This is known as the multimode system. In those cases the more than 1 bit is used by the CPU to set and handle the mode.
- An example of the multimode system can be described by the systems that support virtualization. These CPU's have a separate mode that specifies when the virtual machine manager (VMM) and the virtualisation management software is in control of the system.
- For these systems, the virtual mode has more privileges than user mode but less than kernel mode.

# What is virtualization?


- **Virtualization** -- the abstraction of computer resources.
- Virtualization hides the physical characteristics of computing resources from their users, be they applications, or end users.
- This includes making a single physical resource (such as a server, an operating system, an application, or storage device) appear to function as multiple virtual resources; it can also include making multiple physical resources (such as storage devices or servers) appear as a single virtual resource.



# Virtualization

- Similar to Abstraction but doesn't always hide low layer's details
- Real system is transformed so that it appears to be different
- Virtualization can be applied not only to subsystem, but to an *Entire Machine*  
→ ***Virtual Machine***

# Virtualization



- **Virtualization**, in computing, is the creation of a virtual (rather than actual) version of something, such as a hardware platform, operating system, storage device, or network resources.
- Virtualization is the process by which one computer hosts the appearance of many computers.
- It is used to improve IT throughput and costs by using physical resources as a pool from which virtual resources can be allocated.

# VIRTUALIZATION BENEFITS

- Sharing of resources helps cost reduction
- Isolation: Virtual machines are isolated from each other as if they are physically separated
- Encapsulation: Virtual machines encapsulate a complete computing environment
- Hardware Independence: Virtual machines run independently of underlying hardware
- Portability: Virtual machines can be migrated between different hosts.
- Cross platform compatibility
- Increase Security
- Enhance Performance
- Simplify software migration

# What is “*Machine*”?

- 2 perspectives
- From the perspective of a process
  - ABI provides interface between process and machine
- From the perspective of a system
  - Underlying hardware itself is a machine.
  - ISA provides interface between system and machine

# System/Process Virtual Machines

Can view virtual machine as:

- Process virtual machine
  - Virtual machines can be instantiated for a single program (i.e. similar to Java)
  - Virtual machine terminates when process terminates.
- System virtual machine (i.e. similar to cygwin)
  - Full execution environment that can support multiple processes
  - Support I/O devices
  - Support GUI

# Virtual Machines

- A *virtual machine* takes the layered approach to its logical conclusion. It treats hardware and the operating system kernel as though they were all hardware.
- A virtual machine provides an interface *identical* to the underlying bare hardware.
- The operating system creates the illusion of multiple processes, each executing on its own processor with its own (virtual) memory.

- The resources of the physical computer are shared to create the virtual machines.
  - CPU scheduling can create the appearance that users have their own processor.
  - Spooling (simultaneous peripheral operations online) and a file system can provide virtual card readers and virtual line printers.
  - A normal user time-sharing terminal serves as the virtual machine operator's console.




# VM types

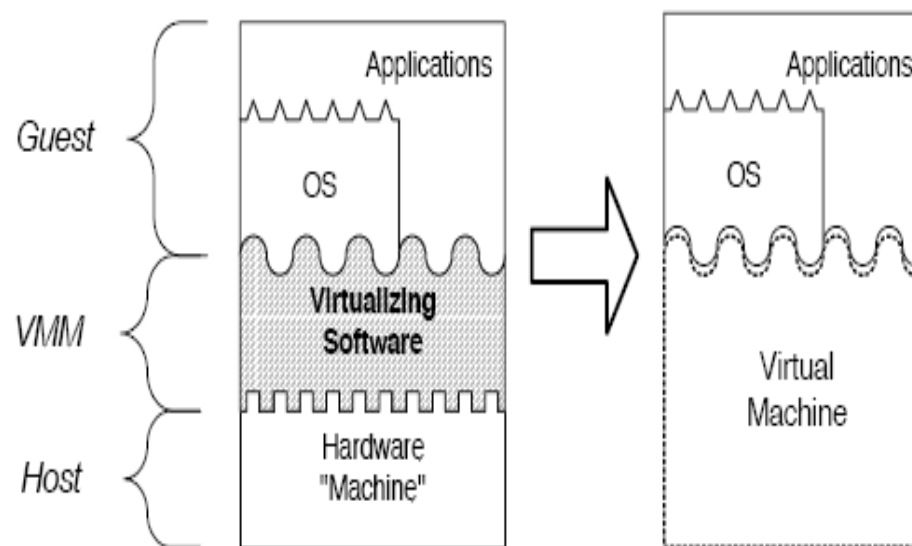
- Native VM
- Hosted VM
- Dual mode VM

# Native VM

- The VM can be provisioned to any hardware system.
- Virtual software placed between underlying machine and conventional software
  - Conventional software sees different ISA from the one supported by the hardware
- The VM is built with virtual resources managed by a guest OS to run a specific application. Between the VMs and the host platform, we need to deploy a middleware layer called a *Virtual Machine Monitor (VMM)* .



**(b) Native VM**




- Figure b shows a native VM installed with the use a VMM called a *hypervisor at the privileged mode*.
- For example, the hardware has a x-86 architecture running the Windows system. The guest OS could be a Linux system and the hypervisor is the XEN system developed at Cambridge University.
- This hypervisor approach is also called bare-metal VM, because the hypervisor handles the bare hardware (CPU, memory, and I/O) directly.

- Virtualization process involves:
  - Mapping of virtual resources (registers and memory) to real hardware resources
  - Using real machine instructions to carry out the actions specified by the virtual machine instructions

# Hosted VM


Another architecture is the host VM shown in Fig. (c). Here the VMM runs with a non-privileged mode. The host OS need not be modified.



(c) Hosted VM

# Dual Mode VM

The VM can be also implemented with a dual mode as shown in Fig. (d).




(d) Dual-mode VM

- Part of VMM runs at the user level and another portion runs at the supervisor level. In this case, the host OS may have to be modified to some extent.
- Multiple VMs can be ported to one given hardware system, *to support the virtualization process.*
- *The VM approach offers hardware-independence of the OS and applications.*
- The user application and its dedicated OS could be bundled together as a virtual appliance, that can be easily ported on various hardware platforms.



# System Models



# Advantages/Disadvantages of Virtual Machines

- The virtual-machine concept provides complete protection of system resources since each virtual machine is isolated from all other virtual machines. This isolation, however, permits no direct sharing of resources.
- A virtual-machine system is a perfect vehicle for operating-systems research and development. System development is done on the virtual machine, instead of on a physical machine and so does not disrupt normal system operation.
- The virtual machine concept is difficult to implement due to the effort required to provide an *exact* duplicate to the underlying machine.